



UNIVERSITÀ DEGLI STUDI DI MILANO - BICOCCA

Scuola di Scienze

Dipartimento di Informatica, Sistemistica e Comunicazione

Corso di Laurea Magistrale in Informatica

Data Analytics Network and Sentiment Analysis on Italian Amazon Reviews of Books

Cocca Umberto - 807191

Academic Year 2020 - 2021

Indice

Introduction	2
Data Analysis Methodologies	3
Network Analysis	3
Sentiment Analysis	3
Existing Software and Tools used	4
Python	4
ASUM	4
Dataset	6
Network Analysis	8
Page Rank	9
Aspect-based sentiment analysis	11
Preprocessing	11
Most common words	14
Sentiword seeds	15
ASUM	16
Conclusions	19

Introduction

In the last years more and more researches have broadened the understanding of textual resources, leading to the growth of online services that changed the face of shopping.

E-commerce applications like Amazon acquire a disproportionate amount of data through their transactions and users, a substantial part is indeed given by the contents generated by users who evaluate the products purchased and share their experience with numerical evaluations and/or reviews.

Sentiment Analysis extracts structured data from textual resources, allowing a statistical analysis of buyer community trends under different aspects. Companies always want to discover opinions and emotions of their clients about their products and services. In addition, even potential costumers want to know the opinions and experience of users who have already used a certain service or purchased a specific product. Knowing the more or less appreciated elements of a product, according to the different categories of users, allows companies to conduct a better forecast of the market and therefore to implement more efficient business strategies.

Reviews are perceived as reliable sources, thus representing a very powerful tool.

It is possible to find the code of this project in this **repository**.

Data Analysis Methodologies

Network Analysis

The goal of this project was to extract insights that may turn helpful for business purposes. In particular, the question I want to answer by using network analysis is: *Which are the most recommended books?* This can be useful to understand how to sort the products, for example within a website, in order to show to users first the ones they are most likely looking for.

Sentiment Analysis

As said before, retrieving customer feedback is an extremely important operation for companies, especially for those that do not mind spending millions for this service. For a long time this has been done manually, and still it is in some cases. In the present days, however, given the enormous amount of information (e.g, clients' feedbacks) it is fundamental to structure the data collection process and the analysis for the decision-making process. For products and services opinions it is possible to use Sentiment Analysis, that can provide answers regarding the most important questions from the customers' point of view.

Sentiment Analysis is used to interpret natural language and identify subjective information that denote opinions, emotions and feelings, determining the corresponding polarity (positive, negative or neutral) and finally summarizing this data so that it can be of value for a company. In this way, decisions can be made based on meaningful data rather than from simple intuitions that are not always correct. Sentiment Analysis is important because companies want their brand to be positively perceived. In this regard, the focus can be on positive or negative comments, as well as customers' feedbacks, to evaluate both strenghts and point on which to improve.

In order to apply Sentiment Analysis in this project, first the textual parts of the reviews are systematically analyzed to extract an opinion. A preliminary pre-processing phase will prepare the dataset and finally, ASUM (Aspect Sentiment Unification Model) is used

to extract set of topics that refer to positive and negative sentiments from a document made of sentences.

Existing Software and Tools used

Python

For the preprocessing and network analysis I used Python, due to the large amount of open source tools and libraries available. In particular, the following libraries were used:

- *Pandas*: to load and manipulate the dataset;
- *igraph*: is a collection of network analysis tools with the emphasis on efficiency, portability and ease of use;
- *NLTK*: to split every review in a list of sentences;
- *re*: to perform a partial cleaning of the data , for example deleting words composed by inadequate characters.

ASUM

Using Python, the ad-hoc input for the Java version of ASUM was built. ASUM was created by Yohan Jo and Alice Oh, it is available at the following [link](#).

The program input consists of two mandatory files and an optional one:

- *BagOfSentences.txt* (mandatory)
This file is a representation of the word list of documents in the corpus. For each document, the first line is the number of sentences, from the next line and on there is a list of indexes that refer to the relative position of a word in the WordList file.
- *WordList.txt* (mandatory)
The file maps words with indexes. It is assumed that the first word has index 0, the second has index 1 and so on.
- *SentiWords-0.txt*, *SentiWords-1.txt*, ... (optional)

These files are composed of words called "semi-sentimental". The files enumeration should start from 0 and then gradually increase, until the number of searched sentiments is reached. In the ASUM model it is possible to help the sampling process by making use of this a priori information. If, for example, we know that a given word is positive because it belongs to the lexicon of positives, then its probability of being positive is known.

For this project two sentiments were searched, one positive and one negative.

Dataset

The dataset is in JSON format and it is loaded into memory as DataFrame, through the Pandas library. Since loading and pre-processing of dataset are the most challenging operations, computationally speaking, I have applied the *.to_pickle()* function to serialize and store the created dataframe in memory in order to speed up all the process.

The original dataset contains several products extracted from Amazon IT and belonging to different categories, such as beauty, books, video games and electronics, with the corresponding reviews stored in another dataset. Since these products belong to the Italian Amazon market, most of the descriptions and reviews are written in Italian. Each record in the reviews dataset represents a single review made by a user for a certain product on the date indicated. Overall, we have 20460 products and 1988855 reviews.

Each product is described by the following fields:

- `_id`: the unique identifier of the product;
- `title`: the title of the product;
- `category`: the category of the product;
- `avg_rating`: the average rating of its reviews measured in stars, that can go from 1 to 5;
- `reviews_number`: the number of reviews for the product;
- `question_number`: the number of questions made for the product;
- `pictures`: the links to the product's images;
- `description`: the description of the product;
- `features`: the list of characteristics of the product;
- `versions`: the list of the other versions of the same product;
- `bought_together`: the list of products often bought together with this product;
- `also_bought`: the list of products often bought by users who bought this product;
- `also_viewed`: the list of products often viewed by users who viewed this product.

Instead, each review is described with the following fields:

- `_id`: the unique identifier of the review;
- `product`: the identifier of the product which the review refers to;
- `title`: the title of the review;
- `author-id`: the unique identifier of the author of the review;
- `author-name`: the name of the author of the review;
- `date`: the date of the review;
- `rating`: the rating of the review measured in stars, that can go from 1 to 5;
- `helpful`: the number of users who rated the review useful;
- `verified`: field that is true if the user actually bought the product the review refers to, false otherwise;
- `body`: the content of the review.

Since the dataset contains a large variety of products, I made the decision to focus just on the books category, reducing the dataset dimension to 558 products and 40097 reviews.

Network Analysis

In order to discover which are the most recommended books, I converted the dataset to a graph with the *igraph* module, where the nodes represent the products and the edges the recommendation relationships between products.

Statistics of the graph	
Nodes	20459
Edges	61393
Directed	True
Density	0.000146
Reciprocity	0.565056
Assortativity	0.067352
Average Degree	34.56

I extracted the relationships from the *bought_together*, *also_bought* and *also_viewed* fields of the dataset and weighted them according to their importance. In particular, the weights of the edges have been assigned as follows:

- 0.5 if the products were linked together in a *bought_together* relationship;
- 0.3 if the products were linked together in a *also_bought* relationship;
- 0.2 if the products were linked together in a *also_viewed* relationship.

Since the relationships were not mutually exclusive (that is, a product could be recommended by another product considering more than one kind of relationship), I weighted the edges connecting two products considering all the possible relationships between them.

Note that the minimum weight for an edge was 0.2, while the maximum weight was 1. The greater the weight of the edge, the stronger is the relationship. I chose these weights because, intuitively, two products frequently bought together have a stronger recommendation relationship than two products frequently viewed together.

Using the indegree function I got the probability distribution that allows to quantify the importance of each book. With this measure, however, I didn't take into account the quality of each of these links. In fact, a link coming from *w* might not meant as much

as a link coming from z , because z is more popular than w . Below the top 15 results of indegree measure, which takes into account the entire network but it is restricted to output only books.

Indegree List	
Name	Id
La versione di Fenoglio	8806240986
#NONOSTANTE	8891821659
A un metro da te	8804709367
Io, te e il mare	8804687290
Oxford Advanced Learner's Dictionary...	0194798798
Dizionario francese...	8867314386
I colori delle emozioni...	8858017390
Grammatica attiva della lingua tedesca...	8820351854
Grammar and vocabulary	1107481112
Un cuore in mille pezzi...	8868363844
Pensa e arricchisci te stesso	8871527151
Padre ricco padre povero...	8871527747
Le parole di Sara	8817109924
1984	8804668237
Le NOSTRE emozioni	8891820024

Page Rank

I computed the Page Rank centrality for every node using the Page Rank algorithm, which is designed to find the most recommended nodes in a network. The weights on the edges have been taken into account because they represent the strength of the recommendation relationship. The top 15 results are shown in Page Rank List, that takes into account the entire network.

Page Rank List	
Name	Id
La versione di Fenoglio	8806240986
Le parole di Sara	8817109924
Km 123	8804716371
#NONOSTANTE	8891821659
Pensa e arricchisci te stesso	8871527151
A un metro da te	8804709367
Io, te e il mare	8804687290
Un cuore in mille pezzi. After: 2	8868363844
Come mondi lontani. After: 3	8868363852
Grammatica attiva della lingua tedesca...	8820351854
Padre ricco padre povero. Quello che...	8871527747
Anime perdute. After: 4	8868363860
Una mutevole verità	8806226487
L'estate fredda	8806227742
#VALESPO	889182254X

Aspect-based sentiment analysis

The main question I want to answer is: what users think about “Il Piccolo Principe”? Since there are 3 different editions of this book, I plotted for each one the number of reviews, in order to focus on the most reviewed one.

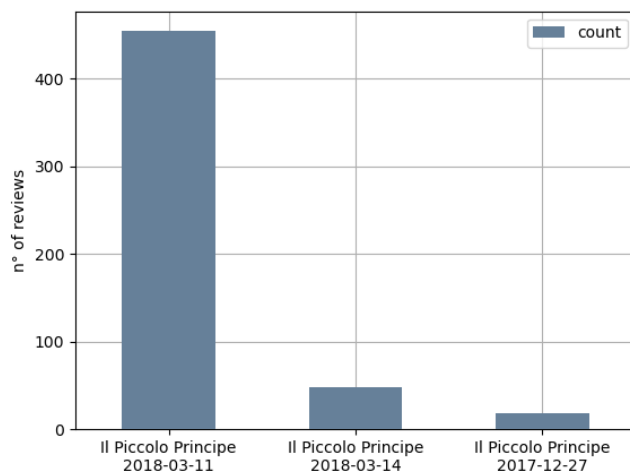


Fig. 1: *n° of review for each "Il Piccolo Principe" edition*

I accomplished this task by using ASUM, a model to perform aspect-based sentiment analysis, with the objective of discovering both the positive and negative aspects. To carry out this study I focused my attention only on the reviews of the most reviewed edition of the book “Il Piccolo Principe”, so I went from the original 36035 book reviews to 454. The most important information of the dataset are the title and the body of the reviews. In particular, in this phase I merged these two columns into one, because the title contains a lot of useful information about the review. Moreover, I selected only the verified reviews (those written by people who actually bought the book) and they represent the 83% of the total reviews, thus 379.

Preprocessing

Before diving into the aspect based sentiment analysis, a preprocessing phase was required. First of all, the fields deemed superfluous for analysis have been removed from the dataset.

The manipulation of the data was done sequentially and with standard steps for this kind of analyzes. For each sentence these steps of preprocessing were applied:



Fig. 2: *Preprocessing phases*

The first phase was tokenization: each of the 379 reviews was tokenized into a list of phrases through the NLTK python library.

In first instance, since they can reverse the sentiment of the sentence, I decided to handle the negations simply placing the `not_` tag in front of all the words that followed the word “non” or the first adversative copulative word up to the first punctuation mark, such as: `[“non”, “ma”, “però”, “invece”, “anzi”, “bensì”, “tuttavia”, “nondimeno”, “pure”, “eppure”]`

After seeing the results I thought this approach was not the best, because they were almost meaningless. Thus, I reduced the list of the negated words to only the “non”, since especially for the other words it’s very common in italian to have these followed by positive words; negating those i would get a completely negative review while originally the first half of the sentence was negative but the second half was positive. Nevertheless, also in this case I saw strange cases of nouns or verbs that should be complementary but were inside the same aspect (e.g “libro” and “not_libro”). For this reason I decided to apply the negation only to adjectives, that led in better final results.

As regards the lemmatization phase, I tried first to apply stemming without lemmatization, but this resulted very inaccurate because Italian verbs have different tenses and, thus, many different tokens related to the same verb were produced. To obtain good re-

sults I standardized them to their infinitive form using a POS tagging called TreeTagger, that it is also able to do grammatical analysis.

Through this TreeTagger I also handled the negations for each adjective that follows the "non" word in a sentence up to the first punctuation mark.

Finally, each sentence was tokenized to create a list of words. All the words that belong to punctuation, strange symbols and stopwords (such as articles or conjunctions, that do not actually help you understand what someone is talking about) were deleted.

As I scrolled through the words, a dictionary WordList.txt (necessary to run ASUM) was built. In parallel I created an indexing, such that each word in the sentences was mapped to its corresponding word from the dictionary to generate the file BagOfSentences.txt, mandatory for ASUM as input. Below is an example of document construction:

```
6
0 1 2 3 4 5 6 7 8 9
10 11 12 13 14
15 16 17 18 19 20 4 21 22 23 24 25 26 27 28 29 30 31 32
33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48
4 40 33 49 34 50 51 52
53 54 55 56 57 4 58
4
59 60 4 40 61 49 62 63 4 64 65 66 67 68 69 70 28 71 72 73 74 75 76 77 40 78 79 80 80 81 82
83 14 84 85 86 87 88 89 90
91 92 1 93 94 79 95 96 97 98 99 98 75 100 101
102 15 103 104 105 106
```

Fig. 3: *esempio input BagOfSentences.txt*

The first line indicates the number of phrases in the review, then there are a number of lines in which each index refers to the position of the word in the wordList.

According to the authors of ASUM, the model works better if the sentences within a review are composed of many words, because short sentences lack of evidence for both sentiments and aspects, and I also removed sentences that were too long. Thus, for the analysis I removed the sentences with:

- *long sentences*: more than 300 words;
- *short sentences*: less than 5 words.

Before starting the experimentation with ASUM, I computed the most common words with respect to positive and negative reviews with the purpose of discovering the most used one. I divided the reviews in positive and negative, considering the number of stars: a review is positive if its number of stars is greater than 3, negative otherwise. Fig. 4 shows the most common words used in the negative reviews, while Fig. 5 shows the most common words used in the positive ones. As we can see, negative reviewers comment more the shipment and the quality of the cover, while the positive ones are more in general about the book contents and the price.



And these negative words:

costoso	brutto	ritardo	rovinato	rotto
orribile	orrendo	noioso	schifo	ripetitivo
vomitare	problema	inutile	ridicolo	senza
danneggiato	sporcizia	sporco	deformare	deformato
vecchio	male	strappato	evitare	danno
difetto	difettoso	infantile	confuso	banale
impossibile	negativo	puerto	confusionario	amaro
noia	difficile	illeggibile	not_adorare	not_adorato
not_bella	not_bellezza	not_bellissima	not_bellissimo	not_bello
not_ben	not_bene	not_allegria	not_allegro	not_amicizia
not_apprezza	not_apprezzabile	not_apprezzamento	not_apprezzare	not_apprezzato
not_appropriato	not_felicità	not_fedeltà	not_felice	not_felicemente
not_felici	not_felicissimo	not_felicitarsi	not_geniale	not_consiglio
not_consigliato	not_gioia	not_lodevole	not_splendidamente	not_splendido
not_squisito	not_successo	not_top	not_vero	not_apprezzabile

ASUM

The ASUM output is composed by some CSVs, I focused on those related to probabilities.

The ASUM parameters are:

- t : number of topics (aspects);
- s : number of sentiments;
- d : number of seed words;
- α : dirichlet prior on the document-sentiment-topic distribution;
- β : dirichlet prior on the word-sentiment-topic distribution;
- γ : dirichlet prior document-sentiment distribution;
- i : number of sampling iterations.

I decided to set $s = 2$ because I was interested in finding only positive and negative

aspects. Since about 87% of the reviews was positive, I set γ equal to 0.9/0.1, where 0.9 is associated with the positive sentiment and 0.1 with the negative sentiment. Then I executed ASUM with different combinations of the remaining parameters α , β and γ . For each execution I manually inspected and interpreted the results. I tried with the settings $t = 3, 5$ and 10.

The general observed behavior is that by increasing the parameter t , it is possible to visualize a series of predominant words that do not appear when using lower t , moreover the aspects are more defined (e.g. with $t = 3$ words related to both shipment and edition where under the same topic). In fact, increasing t we have positive aspects more connected to different emotions. The problem, however, is that as seen previously during the preprocessing phase the results of the negative aspects had almost no meaning, the reason of this may be linked to the actual number of negative topics, problems with the handling of negation or with the lexicon.

Increasing the number of topics up to 10 I found some duplicates within the positive aspects, actually I saw duplicates regarding the emotions of the users. With negative aspects, however, some aspects weren't truly negative, but with the right workflow they were more defined. Despite the duplicates, the results were no doubt more interesting. I did not notice any relevant differences in the results by varying of the values for α and β .

In conclusion, I isolated 5 negative and 4 positive aspects. The most relevant topics, both positive and negative, are reported in the following tables:

Positive				
Edition	Shipment	Classic book	Book quality	Memories
libro	arrivar	libro	libro	uomo
copertina	ottimo	bambino	ottimo	ricordare
rigido	buono	lettura	carattere	infanzia
ottimo	amazon	leggero	leggibile	bambino
copia	spedizione	classico	qualità	ricordo
edizione	veloce	adulto	stampa	not_triste
economico	servizio	consigliare	lettura	racconto
stampato	consegna	acquistare	dimensione	storia
pagina	impeccabile	rileggere	pagina	adulto
rilegatura	ordinare	regalare	carta	favola

Table 1: Positive aspects

Negative			
Cost	Shipment	Book quality	Book content
libro	ridicolo	libro	male
costo	rimborso	problema	leggero
spendere	amazon	traduzione	inutile
inutile	problema	pagina	ridicolo
traduzione	solo	errore	noia
qualità	tardi	stampa	pensare
leggero	spedizione	rilegare	prezzo
noioso	tutelare	pessimo	vita
not_apprezzabile	danno	incollare	solo

Table 2: Negative aspects

Finally, I have verified the results obtained by ASUM by filtering the reviews using the aspects as keywords and by manually inspecting them.

Conclusions

In this report I extracted valuable insights from books using a dataset containing products and reviews from the Italian Amazon market. The first goal of this project was to discover the most recommended product of the category "book" in the Amazon network and then figure out what people think about one of these, the book "Il Piccolo Principe" by Antoine de Saint-Exupéry. Concerning the first goal, I discovered that the most recommended book on Amazon is "La versione di Fenoglio". Then, I discovered that most people are very satisfied about "Il Piccolo Principe", regarding the edition, the quality of the book and it is overall considered a classic book, recommended both for adults and children, especially as a gift. There are however conflicting opinions about the shipment and the book's print quality. Some people also didn't like the book content and were not satisfied with the price.